

C++ for Numerical Programming - lectures 1-3

Martin Robinson

2017

A few introductory remarks

- C++ is a common programming language for high performance scientific computing applications, but it is not the only choice, know what languages/platforms are available to you.
- It is generally useful to know at least one low-level language (e.g. C++) and one high-level language (e.g. Python, Matlab) for your work.
- C++ is useful for:
 - High performance applications. Concurrency, parallelism and compilation to efficient, optimised binaries is easier in C++
 - Software libraries. C++ is designed for zero-cost abstraction, so library authors can provide their functionality with minimal overhead.
 - Access to a truly vast array of third-party scientific libraries

- There are lots of resources for C++, particularly online.
 - Many books, e.g. Bjarne Stroustrup, “C++ Programming Language”. Make sure to use editions published well after 2011.
 - Reference material: <http://www.cplusplus.com> and <http://cppreference.com>
 - Q&A sites, <http://stackoverflow.com> etc.
- Integrated Development Environments (IDEs) exist—e.g. Eclipse or CLion. These are useful for beginners as well as managing large code development. A command-line editor such as `vim` can also be highly customised for C++ development

Programming paradigms in C++

C++ is a multi-paradigm language with the goal of allowing powerful abstractions with *zero* run-time cost

- *Procedural*: The C-style basis of C++, algorithm consists of procedures (i.e. functions or subroutines) containing a series of computational steps
- *Object-orientated*: break computational task into objects (i.e. classes) that expose behaviours (methods) and data (members) through a public interface
- *Statically typed*: Variables have a constant type determined at compile-time, rather than determined on-the-fly at run-type (e.g. MATLAB)

Programming paradigms in C++

Recent updates (e.g. C++11) focus on generic and functional paradigms

- *Generic*: algorithms are written with generic types or variables to be determined later (at compile-time). This is achieved in C++ through *templates*
- *Functional*: Not a *pure* functional language, but functional elements exist in C++, for example the Standard Template Library (STL) and lambda functions.

This course gives you a practical toolbox of C++ programming up to C++11. *This is a small part of C++ as a whole.*

Lecture 1 — The basics

General structure of a basic C++ program

```
#include <header1>
#include <header2>
int main(int argc, char* argv[])
{
    line of code;
    // this is a comment, ignored by the compiler
    more code; // this is a comment as well
    /*
        * Multi-line
        * comment
        */
    return 0;
}
```

Things to note

- ① Header files are listed first. These are files that contain the functions needed for operations such as input, output and mathematical calculations
- ② There is a section of code that starts `int main(int argc, char* argv[])` (or `int main()` for brevity when command-line input is unimportant)
- ③ This section of code is followed by more code enclosed between curly brackets, `{` and `}`
- ④ Comments may easily be inserted into the code
- ⑤ Lines of code that “do something” end with a semicolon `;`
- ⑥ Just before the closing curly bracket at the end of the code is a statement `return 0;`

A first C++ program

```
#include <iostream>
int main(int argc, char* argv[])
{
    std::cout << "Hello World\n";
    return 0;
}
```

Save this code as `hello_world.cpp`

This program prints the text “Hello World” to the screen

- `iostream` is a header file that is needed when using input and output
- `std::cout` is a command that sends output to the console, i.e. the screen
- `\n` is a formatting command that starts a new line
- All statements (lines of the program) inside the curly brackets end with a semicolon ;

Compiling the code

A key difference between Matlab and C++—before the code can be executed it must be compiled. When using the Gnu compiler of Unix/Cygwin this code can be compiled by saving the code and typing

```
$ g++ -Wall -O -o hello_world hello_world.cpp
```

followed by return. This produces an executable called hello world that can be executed by typing

```
$ ./hello_world
```

Numerical variables

Before a variable is used the type of variable must be declared. For example if the variables `i` and `j` are integers and `a` is a double precision floating point number the statements

```
int i, j;  
double a;
```

must be included in the program before these variables are used. It is advisable to use `double` rather than `float` in scientific computing applications

Some names, such as `int`, `for`, `return` may not be used as variable names because they are used by the language. These words are known as reserved words or keywords

Some example lines of code that add two integers and print the answer to screen are

```
int integer1, integer2, answer;  
integer1 = 5;  
integer2 = 10;  
answer = integer1 + integer2;  
std::cout << "The sum of " << integer1 << " and "  
    << integer2 << " is " << answer << "\n";
```

The details of the `std::cout` statement will be explained later

More on numerical variables

A variable may be initialised when defining the variable type, for example

```
int i = 5;
```

The value of more than one variable may be assigned in each statement

```
int i, j;  
i = j = 3;
```

The second statement in the box above may cause confusion—it actually means

```
i = ( j = 3 );
```

and not

```
( i = j ) = 3;
```

in which *j* is unchanged

When assigning values to floating point variables it is good programming practice to write numbers with decimal points, i.e.

```
double a = 5.0;
```

rather than

```
double a = 5;
```

If a quantity is a constant throughout the program it may be declared as such

```
const double density = 45.621;
```

There is a shorthand for some mathematical operations

Longhand	Shorthand
-----	-----
<code>a = a + b;</code>	<code>a += b;</code>
<code>a = a - b;</code>	<code>a -= b;</code>
<code>a = a * b;</code>	<code>a *= b;</code>
<code>a = a / b;</code>	<code>a /= b;</code>
<code>a = a + 1;</code>	<code>a++;</code> if <code>a</code> is an integer
<code>a = a - 1;</code>	<code>a--;</code> if <code>a</code> is an integer

Some variable types and ranges in most 32-bit operating system are given below. Note the ranges are operating system dependent

Variable type	C++ name	Range in 32-bit
-----	-----	-----
integer	int	-2^{31} to $2^{31} - 1$
integer	short int	-2^{15} to $2^{15} - 1$
integer	long int	-2^{31} to $2^{31} - 1$
unsigned integer	unsigned int	0 to $2^{32} - 1$
floating point	float	-3.4×10^{38} to 3.4×10^{38}
floating point	double	-1.7×10^{308} to 1.7×10^{308}
floating point	long double	-1.7×10^{308} to 1.7×10^{308}

Note that double and long double are identical

Note that the size of long int varies with 64-bit architectures

Some example lines of code

```
float a, b;  
long double d, e;  
a = 3.0;  
b = ( a * std::pow(a, 3) ) / 2.0;  
d = 4.0;  
e = 2.0 * std::sqrt(d);
```

`std::pow(x,y)` gives the value of x^y . The `std::` indicates the standard namespace.

`std::sqrt(d)` gives the square root of the variable `d`

When using mathematical functions such as `std::pow`, you should use the additional header file `cmath` and so you need the line of code

```
#include <cmath>
```

Division of an integer by another integer will almost certainly cause problems

An example is given in the following piece of code

```
int i = 5, j = 2, k;  
k = i / j;  
std::cout << k << "\n";
```

The variable `k` is an integer and so cannot store the true value, 2.5
Instead, it will store the value 2

Suppose an integer is divided by a variable of type `long double` - or vice versa - and that the result returned is stored in a variable of type `long double`, as shown in the code below. The variable `k` in this code will contain the mathematically correct answer

```
long double i = 5.0, k;  
int j = 2;  
k = i / j;  
std::cout << k << "\n";
```

However, in general you should construct operations on variables of the same type, for clarity.

Suppose an integer is divided by a variable of type `long double` - or vice versa - and that the result returned is stored in a variable of type `int`, as shown in the code below.

The variable `k` in this code is unable to store the mathematically correct answer

```
long double i = 5.0;
int j = 2, k;
k = i / j;
std::cout << k << "\n";
```

Only construct mathematical operations that are on elements of the same type

An integer can be converted, or cast, to a different data type – see the next slide

Variables can be converted from one type to another for example

```
long double i = 5.0, k;  
int j = 2;  
k = i / static_cast<long double>(j);  
std::cout << k << "\n";
```

In this example, `((long double) (j))` allows the variable `j` to behave as if it were a `long double` variable.

ASCII characters

ASCII characters are numbers, uppercase letters, lowercase letters and some other symbols

These characters may be represented using the data type `char`

```
#include <iostream>
int main()
{
    char letter;
    letter = 'a'; // note the single quotation marks

    std::cout << "The character is " << letter << "\n";

    return 0;
}
```

Boolean variables

These variables take the values `true` or `false`, and are of use when using `if` conditionals and `while` loops

They are used as follows:

```
bool flag;  
flag = true;
```

Strings

A character is one letter or number, a string is an ordered collection of characters

For example, "C++" is a string consisting of the ordered list of characters "C", "+", and "+"

To use strings in C++ requires an extra header file as shown below

```
#include <iostream>
#include <string>
int main()
{
    std::string city; // note the std::
    city = "Oxford"; // note the double quotation marks
    std::cout << city << "\n";
}
```


The if statement

Suppose you want to execute two statements if the condition $p > q$ is met

This is achieved using the following code

```
if (p > q)
{
    statement1;
    statement2;
}
```

Note the indentation of the loop. This makes it clear which statements are executed in the loop

If only one statement is to be executed curly brackets aren't strictly necessary

For example, the following code will execute statement1 if the condition $p > q$ is met

```
if (p > q)
    statement1;
```

but this is poor software engineering practice

Instead, write this code as

```
if (p > q)
{
    statement1;
}
```

The use of curly brackets makes it clear which statement(s) are to be executed

Third example – more than one condition

```
if (p == 0)
{
    statement1;
}
else if (p < 0)
{
    statement2;
    statement3;
}
else
{
    // p > 0
    statement4;
}
```

Fourth example – nested if statements

```
if (p < q)
{
    if (x >= y)
    {
        statement1;
    }
}
```

Fifth example – more than one condition

```
if (p < q || x < y)
{
    statement1;
}
```

statement1 is executed if and only if one or both of $p < q$ and $x < y$ is true - i.e. `||` is the logical OR operator

Relational and logical operators

relation	operator
-----	-----
equal to	== (note that it isn't "=")
not equal to	!=
greater than	>
less than	<
greater than or equal to	>=
less than or equal to	<=

logical condition operator

AND	&&
OR	
NOT	!

Boolean variables may be used in if statements as follows

```
bool flag1 = true, flag2 = false;
if (flag1)
{
    std::cout << "Does print something" << "\n";
}
if (flag2)
{
    std::cout << "Doesn't print anything" << "\n";
}
if (!flag2)
{
    std::cout << "Does print something" << "\n";
}
```

The while statement

Similar syntax to if statements

```
while ( x < 100.0 && i < 10 )  
{  
    x += x;  
    i++;  
}
```

The condition $x < 100.0 \ \&\& \ i < 10$ is tested only at the beginning of the statements in the loop, and not after every statement.

For example if the loop is entered when $x = 99.0$ and $i = 1$, the loop will be executed completely once.

Loop won't be entered when $x \geq 100$: x and i will be unchanged. If you need a loop to execute at least once, with a test at the end, use `do { ... } while (condition)`

for loops

The following loop executes the statements inside the loop 10 times. Note that `i` must have been declared as an integer earlier in the code

```
for (i=0; i<10; i++)  
{  
    statement1;  
    statement2;  
}
```

Alternatively, the integer `i` may be declared inside the loop:

```
for (int i=0; i<10; i++)  
{  
    statement1;  
    statement2;  
}
```

for loops can be nested and run over variable indices. The output from the following section of code is given on the next slide

```
for (int i=0; i<5; i++)  
{  
    for (int j=5; j>i; j--)  
    {  
        std::cout << "i = " << i << " j = " << j << "\n";  
    }  
}
```

The output from the section of code on the previous slide is

```
i = 0 j = 5  
i = 0 j = 4  
i = 0 j = 3  
i = 0 j = 2  
i = 0 j = 1  
i = 1 j = 5  
i = 1 j = 4  
i = 1 j = 3  
i = 1 j = 2  
i = 2 j = 5  
i = 2 j = 4  
i = 2 j = 3  
i = 3 j = 5  
i = 3 j = 4  
i = 4 j = 5
```

Random Number generation

You can generate random numbers using the standard library

```
#include <random>
int main()
{
    // generates a pseudo-random number sequence
    std::default_random_engine generator;

    // uses generator to output a uniform distribution
    std::uniform_real_distribution<double> uniform(0,1);

    // print out the sequence
    for (int i = 0; i < 100; ++i)
    {
        std::cout << uniform(generator) << std::endl;
    }
    return 0;
}
```

Tip: Integrated development environments (IDEs)

- This course tries to be IDE agnostic, only assumes an editor and a command-line compiler
- IDEs such as Visual Studio, CLion and Eclipse, provide in-built debuggers, visual warnings, build management, auto-completion etc.

Lecture 3 — Fixed sized arrays, and IO

If the size of the array is known in advance (i.e. at compile-time) then it is better to use the fixed size array `std::array<T,N>`, where `T` is the type that the array will hold (e.g. `int`, `double`), and `N` is the length of the array.

```
#include <array>
std::array<int,4> x; // x hold 4 integers
const size_t size = 5;
std::array<double,size> y; // y holds 5 doubles

// a 5 by 5 array
std::array<std::array<double,size>,size> z;
```

In contrast to Matlab and FORTRAN the indices of an array of length `n` start at 0 and end at `n-1`

Elements of the array are accessed by placing the indices in separate square brackets, for example

```
x[0] = 1;  
z[1][2] = 3.0;
```

Arrays can be initialised when they are declared, for example

```
std::array<double,3> array1 = {0.0, 1.0, 2.0};  
std::array<std::array<double,3>,2> array2 =  
    {{{0.0, 1.0, 2.0},  
      {3.0, 4.0, 5.0}}};
```

where array2 represents the matrix

$$\begin{pmatrix} 0 & 1 & 2 \\ 3 & 4 & 5 \end{pmatrix}$$

Note that the values of arrays may only be set using the curly bracket notation when they are declared - for example the code

```
std::array<int,3> array = {0, 1, 2};
```

is correct, but the code

```
std::array<int,3> array;  
array = {0, 1, 2};
```

is not correct.

Console output

Console output may be achieved by using `std::cout`

We have already seen that the statement

```
std::cout << "Hello World\n";
```

prints the text “Hello World” to the screen, followed by a newline

The statements

```
int x = 1, y = 2;
```

```
std::cout << "x = " << x << " and y = " << y << "\n";
```

give output to the screen

x = 1 and y = 2

Note that any spaces required in the output must be included within quotation marks

Some useful formatting commands are

Command	Symbol
-----	-----
newline	\n
tab	\t
,	\,
"	\"
?	\?
(bell)	\a

Sometimes, for example if the computer is busy doing a large volume of computation, the program may not print the output to the screen immediately. If immediate output is desirable then use `std::cout.flush()` after any `std::cout` commands

```
std::cout << "Hello World\n";  
std::cout.flush();
```

or use `std::endl`, which combines a newline with a flush

```
std::cout << "Hello World" << std::endl;
```

Keyboard input

Keyboard input for numbers and characters is achieved using

`std::cin`

The following statements prompts someone to enter their PIN

```
int pin;
std::cout << "Enter your PIN, then hit RETURN\n";
std::cin >> pin;
```

`cin` may be used to ask for more than one input at a time

```
int accountno, pin;
std::cout << "Enter your account number then hit RETURN,\n";
std::cout << "and then your PIN followed by RETURN\n";
std::cin >> accountno >> pin;
```

Keyboard input for strings is slightly different. An example is given below

```
#include <iostream>
#include <string>
int main()
{
    std::string name;
    std::cout << "Enter your name and then hit RETURN\n";
    std::getline(std::cin, name);
    std::cout << "Your name is " << name << "\n";
    return 0;
}
```

Writing to file

When writing to file, an additional header function `fstream` is needed.

```
#include <iostream>
#include <fstream>
#include <string>
#include <cmath>
```

The file `output.dat` may be opened using the statement

```
std::ofstream out("output.dat");
```

We can then write to this file in a similar manner as writing to the screen, with the exception that `cout` is replaced by `out`

There are a number of formatting options provided by C++ The following prints data in scientific format

```
#include <iostream>
#include <fstream>
int main()
{
    double x = -1.0, y = 45.3275893627129, z = 0.00000001;
    std::ofstream out("output.dat");
    assert(out.is_open());
    out.setf(std::ios::scientific|std::ios::showpos);
    out << x << " " << y << " " << z << "\n";
    out.close();
    return 0;
}
```

Suppose the file `numbers.dat` (in my home directory) has 3 columns of numbers. This file can be opened using the following code

```
#include <fstream>
```

```
double x, y, z;
```

```
std::string line;
```

```
std::ifstream input("numbers.dat");
```

```
assert(input.is_open());
```

`input.is_open()` returns true if the file was successfully opened

We can read the file like so:

```
#include <sstream>
```

```
while (std::getline (input,line))  
{  
    std::istringstream s(line);  
    s >> x >> y >> z;  
}  
input.close();
```

`std::getline` gets the next line of the file, returns false if we are at the end of the file

`std::istringstream` converts the `std::string` `line` to a stream (like `std::cin`) which can be used to separate out the three columns

Use of assert statements for debugging

`assert` statements can be used in code to confirm something you expect to be true

For example, you may wish to confirm that a number you are about to take the square root of is non-negative

If the condition is not met, the code aborts giving an error message that explains what went wrong

To use `assert` statements you must use the header file `cassert`

An example of the use of `assert` statements is given on the next slide

```
#include <iostream>
#include <cassert>
#include <cmath>
int main()
{
    double a;
    std::cout << "Enter a non-negative number\n";
    std::cin >> a;
    assert(a >= 0.0);
    std::cout << "The square root of a is " << std::sqrt(a)
    return 0;
}
```

If the code on the previous slide is compiled and run, and the user enters “-5” at the prompt, the code will be terminated and the following error message given:

```
a.out:: program.cpp:10: int main(): Assertion 'a >= 0.0' fa
```

This error message tells us to look at line 10 in `program.cpp`, where the assertion in quotes failed the test